


▼ Purpose:

fhfhs;fh[ifhoishfoisdfhbjbjfpjufpsifspfjspoifjhsdoijoiivhnjsaovh;ovhv fhojf'sj'fj'F JOPJ'FPOJ'FJ'fpjF KJOJFCP'fcj'PFJP'afcj fioi;fjoifAOFJ

Approach:

FNJ;HF;IOFHWFHWOE;IHJFOWJF;OJ WFCNWOIHF;OIWFHJCWHFC J;OIJH;IOWCVFOIWHOJHWVCFH  
HHFC;WOHVCFOWHVOIVHOIWHCVFOWVHOIWW CVJOIVFCHIVHIO WV;OIHVJ;OIVOJHV WHIFW;OFWOIFHOW;NVVOJHV;H  
VCOWIJV;IOVCJ;CVFNJOI

```
from google.colab import drive
drive.mount('/content/drive')
```

 Mounted at /content/drive

```
file_path = '/content/drive/My Drive/scaler_clustering.csv'
```

```
# Import necessary libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import re
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.impute import KNNImputer
from scipy.cluster.hierarchy import dendrogram, linkage
```

```
# Load the dataset
df = pd.read_csv(file_path)
```

```
# Check the first few rows of the dataset
print(df.head())
```

<

	job_position	ctc_updated_year
0	Other	2020
1	FullStack Engineer	2019
2	Backend Engineer	2020
3	Backend Engineer	2019
4	FullStack Engineer	2019

```
# Checking the shape, data types, and missing values
print(df.shape)
```

(205843, 7)

```
# Check for duplicates
duplicates = df.duplicated()
```

```
# Count the number of duplicate rows
num_duplicates = duplicates.sum()
print(f"Number of duplicate rows: {num_duplicates}")
```

Number of duplicate rows: 0

```
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205843 entries, 0 to 205842
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   index                 205843 non-null int64
1   company_hash          205799 non-null object
2   email_hash            205843 non-null object
3   orgyear               205757 non-null float64
4   ctc                   205843 non-null int64
5   job_position          153279 non-null object
6   ctc_updated_year      205843 non-null int64
dtypes: float64(1), int64(3), object(3)
memory usage: 11.0+ MB
None
```

I see that orgyear (which is "employment start year", ) has afew values which are future years and a few values which lead to more than 50 years of experience, which is not plausible. I would like to remove these rows from the dataset.

```
# Set the current year
current_year = 2024
```

```
# Detect and eliminate rows where 'orgyear' is in the future or leads to more than 50 years of experience
# Define a valid range for 'orgyear': should be less than or equal to the current year, and not more than 50 years old
valid_orgyear = (df['orgyear'] <= current_year) & (df['orgyear'] >= current_year - 50)
```

```
# Filter the dataset to keep only the valid records
df_cleaned = df[valid_orgyear]
```

```
# Check the shape of the cleaned dataframe and number of rows removed
print(f"Original dataset size: {df.shape}")
print(f"Cleaned dataset size: {df_cleaned.shape}")
print(f"Number of invalid records removed: {df.shape[0] - df_cleaned.shape[0]}")
```

↗

Original dataset size: (205843, 7)  
Cleaned dataset size: (205665, 7)  
Number of invalid records removed: 178

```
# Get the invalid rows that were removed
invalid_rows = df[~valid_orgyear]
```

```
# Display the invalid rows
print(invalid_rows)
```

↗

	index	company_hash \			
	2211	2211	phrxkv		
	2333	2333	xgmgn ntwyzgrgsxto ucn rna		
	2562	2562	tj		
	3122	3122	ft tdwtr		
	3365	3365	fyxntyvn lq		
	...	...	...		
	196354	197352	vaxnjv mxqrv wvuxnvr		
	198187	199212	xb v onhatzn		
	202210	203276	mqvmtzatq		
	203992	205068	xatv ouvqp ogrhnxgzo ucn rna		
	205435	206515	vhngsqxa		
		email_hash	orgyear	ctc \	
	2211	3394674bb6bb1de6289e931853fa0bd131c811e0054a92...	2031.0	1500000	
	2333	c737ceb66c7f0ce37c2fce087003aa129632a3a2fa4f6c...	NaN	170000	
	2562	25edac17c77f6f0edeafb86f7a7844d96dc899e193c87e...	NaN	860000	
	3122	c402eba160abf4e5b5f72af775fc98dbd346f1a081112e...	NaN	600000	
	3365	38bd913564fa983cd4fb7799e4027d8ed2b0fd6263e15a...	NaN	800000	
	...	...	...	...	
	196354	069308440811d578c817c05392f97e8919baac6aa12aa3...	1.0	2900000	
	198187	9429a19771ae913f169917d380c94f003115aaaf904388...	2025.0	300000	
	202210	d66f939c4318c1958be5bc9e7b70b741aa61be7493ff58...	2028.0	1300000	
	203992	7191da2e57dc0c1301711e889ea72d5cc801e039359b1...	20165.0	850000	
	205435	3fa8de870da01d863abba8eb6a8ae3df1aa18c18946688...	NaN	2400000	
		job_position	ctc_updated_year		
	2211	Backend Engineer	2020		
	2333	Other	2020		
	2562	Data Analyst	2020		

```

3122    Support Engineer    2020
3365           NaN        2021
...
196354    Data Scientist    2019
198187           Other    2021
202210    Backend Engineer    2021
203992           NaN    2019
205435           NaN    2020

```

```
[178 rows x 7 columns]
```

```
print(df_cleaned.isnull().sum())
```

```

↗ index            0
  company_hash      44
  email_hash        0
  orgyear           0
  ctc               0
  job_position      52508
  ctc_updated_year   0
dtype: int64

```

```

# Mode imputation for company_hash
df_cleaned['company_hash'].fillna(df_cleaned['company_hash'].mode()[0], inplace=True)

```

```

↗ <ipython-input-17-f64e7f3aeabb>:2: SettingWithCopyWarning:
  A value is trying to be set on a copy of a slice from a DataFrame

```

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy).

```
df_cleaned['company_hash'].fillna(df_cleaned['company_hash'].mode()[0], inplace=True)
```

```
print(df_cleaned.isnull().sum())
```

```

↗ index            0
  company_hash      0
  email_hash        0
  orgyear           0
  ctc               0
  job_position      52508
  ctc_updated_year   0
dtype: int64

```

```

# Fill missing values (Mean/ KNN Imputation)
# imputer = KNNImputer(n_neighbors=5)

```


```
# df_cleaned[['orgyear']] = imputer.fit_transform(df_cleaned[['orgyear']])
```

```
# print(df_cleaned.isnull().sum())
```

```


# Convert orgyear
df_cleaned['orgyear'] = df_cleaned['orgyear'].astype(int)

```


 <ipython-input-22-40f87f736f36>:2: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy).  
df\_cleaned['orgyear'] = df\_cleaned['orgyear'].astype(int)


```
print(df_cleaned.info())
```

 <class 'pandas.core.frame.DataFrame'>  
Index: 205665 entries, 0 to 205842  
Data columns (total 7 columns):  
#     Column               Non-Null Count     Dtype  
---  -----  
0     index                 205665 non-null   int64  
1     company\_hash         205665 non-null   object  
2     email\_hash            205665 non-null   object  
3     orgyear               205665 non-null   int64  
4     ctc                   205665 non-null   int64  
5     job\_position         153157 non-null   object  
6     ctc\_updated\_year     205665 non-null   int64  
dtypes: int64(4), object(3)  
memory usage: 12.6+ MB  
None

```
print(df_cleaned.isnull().sum())
```


 index                 0  
company\_hash           0  
email\_hash             0  
orgyear                 0  
ctc                     0  
job\_position            52508  
ctc\_updated\_year        0  
dtype: int64

```
# Group by company_hash and fill missing job_position with mode for that company  
df_cleaned['job_position'] = df_cleaned.groupby('company_hash')['job_position'].transform(lambda x: x.fillna(x.mode()[0] if not x.mode().empty else 'Unknown'))
```

 <ipython-input-25-52a7906e5793>:2: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy).  
df\_cleaned['job\_position'] = df\_cleaned.groupby('company\_hash')['job\_position'].transform(lambda x: x.fillna(x.mode()[0] if not x.mode().empty else 'Unknown'))

```
print(df_cleaned.isnull().sum())
```

 index                 0  
company\_hash           0  
email\_hash             0  
orgyear                 0  
ctc                     0  
job\_position            0  
ctc\_updated\_year        0


dtype: int64

#####

df\_cleaned.head()



	index	company_hash	email_hash	orgyear	ctc	job_position	ctc_updated_year
0	0	atrgxnnt xzaxv	6de0a4417d18ab14334c3f43397fc13b30c35149d70c05...	2016	1100000	Other	2020
1	1	qtrxvzwt xzegwgbb rxbxnta	b0aaf1ac138b53cb6e039ba2c3d6604a250d02d5145c10...	2018	449999	FullStack Engineer	2019
2	2	ojzwnvwnxw vx	4860c670bcd48fb96c02a4b0ae3608ae6fdd98176112e9...	2015	2000000	Backend Engineer	2020
3	3	ngpgutaxv	effdede7a2e7c2af664c8a31d9346385016128d66bbc58...	2017	700000	Backend Engineer	2019
4	4	qxen sqghu	6ff54e709262f55cb999a1c1db8436cb2055d8f79ab520...	2017	1400000	FullStack Engineer	2019




# Unique email hashes and frequency of occurrences  
print(df\_cleaned['email\_hash'].value\_counts())



```
email_hash
bbace3cc586400bbc65765bc6a16b77d8913836cfc98b77c05488f02f5714a4b    10
3e5e49daa5527a6d5a33599b238bf9bf31e85b9efa9a94f1c88c5e15a6f31378     9
6842660273f70e9aa239026ba33bfe82275d6ab0d20124021b952b5bc3d07e6c     9
298528ce3160cc761e4dc37a07337ee2e0589df251d73645aae209b010210eee     9
d15041f58bb01c8ee29f72e33b136e26bc32f3169a40b53d75fe7ae9cbb9a551     8
..
6ed7767a6ba36e8ab4f4d2397a4d32f26f34387720645906bf51a05c2152fd56     1
9778d2fa1bbfb721c3e90941cb3474740610d301f2ccf1429f5c6835ae5e27f4     1
9a891d279335db60cd6a45c2243bca2c56f940e31c5a812a6f642ea800832c4b     1
e96207e084f4552ba131598c704d2c5f12373999fc66285f58dea00afb9d333c     1
0bcfc1d05f2e8dc4147743a1313aa70a119b41b30d4a1f7e738a6a87d3712c31     1
Name: count, Length: 153292, dtype: int64
```


# Remove special characters from Company\_hash  
df\_cleaned['company\_hash\_cleaned'] = df\_cleaned['company\_hash'].apply(lambda x: re.sub('[^A-Za-z0-9 ]+', '', str(x)))



```
<ipython-input-30-eb3fda435fdc>:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
df_cleaned['company_hash_cleaned'] = df_cleaned['company_hash'].apply(lambda x: re.sub('[^A-Za-z0-9 ]+', '', str(x)))
```


# Unique "company\_hash\_cleaned" hashes and frequency of occurrences  
print(df\_cleaned['company\_hash\_cleaned'].value\_counts())



```
company_hash_cleaned
nvnv wgzohrnvzwj otqcxwto    8379
xzegojo                    5378
vbvkgz                     3480
zgn vuurxwvmrt vwwghzn     3408
wgszxkvzn                  3238
...
```

```
bvpt owyggr                1
vhngsqxa xzaxv             1
ctavznh td kteg            1
ihxwprgsxw ogenfvqt       1
bvptbjnqxu td vbvkgz      1
Name: count, Length: 37246, dtype: int64
```


```
# Dropping the 'company_hash' column from the DataFrame
df_cleaned.drop(columns=['company_hash'], inplace=True)
```

 <ipython-input-32-d2a89671a684>:2: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame


See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy).

```
df_cleaned.drop(columns=['company_hash'], inplace=True)
```

```
print(df_cleaned.shape)
```

 (205665, 7)


```
# Create new feature 'Years_of_Experience'
current_year = 2024
df_cleaned['Years_of_Experience'] = current_year - df_cleaned['orgyear']
```



 <ipython-input-34-c3a4f380c3bf>:3: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row\_indexer,col\_indexer] = value instead

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returning-a-view-versus-a-copy](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy).

```
df_cleaned['Years_of_Experience'] = current_year - df_cleaned['orgyear']
```

```
df_cleaned.head()
```



	index	email_hash	orgyear	ctc	job_position	ctc_updated_year	company_hash_cleaned	Years_of_Experience	
0	0	6de0a4417d18ab14334c3f43397fc13b30c35149d70c05...	2016	1100000	Other	2020	atrgxnnt xzaxv	8	
1	1	b0aaf1ac138b53cb6e039ba2c3d6604a250d02d5145c10...	2018	449999	FullStack Engineer	2019	qtrxvzwt xzegwgbb rxbxnta	6	
2	2	4860c670bcd48fb96c02a4b0ae3608ae6fdd98176112e9...	2015	2000000	Backend Engineer	2020	ojzwnvwnxw vx	9	
3	3	effdede7a2e7c2af664c8a31d9346385016128d66bbc58...	2017	700000	Backend Engineer	2019	ngpgutaxv	7	
4	4	6ff54e709262f55cb999a1c1db8436cb2055d8f79ab520...	2017	1400000	FullStack Engineer	2019	qxen sqghu	7	

```
from google.colab import files
```

```
# # Export the cleaned DataFrame to Excel
# df_cleaned.to_excel('scaler_clustering_pre_pro_01.xlsx', index=False)
```

```
# # Download the file
# files.download('scaler_clustering_pre_pro_01.xlsx')
```

```
# Five-point summary of CTC
grouped = df_cleaned.groupby(['company_hash_cleaned', 'job_position', 'Years_of_Experience'])['ctc'].agg(['mean', 'median', 'max', 'min', 'count'])

# Merge it back to the original dataframe
df_cleaned = df_cleaned.merge(grouped, how='left', on=['company_hash_cleaned', 'job_position', 'Years_of_Experience'])

# Create flags for designation, class, and tier
df_cleaned['Designation'] = np.where(df_cleaned['ctc'] > df_cleaned['mean'], 1, 0)
df_cleaned['Class'] = pd.qcut(df_cleaned['ctc'], 3, labels=[3, 2, 1])
df_cleaned['Tier'] = pd.qcut(df_cleaned['ctc'], 3, labels=[3, 2, 1])
```

```
# Export the cleaned DataFrame to Excel
df_cleaned.to_excel('scaler_clustering_pre_pro_02.xlsx', index=False)
```

```
# Download the file
files.download('scaler_clustering_pre_pro_02.xlsx')
```



```
df1=df_cleaned.copy() # used for downsizing for hierarchical clustering at later steps
```

```
df_cleaned.head()
```



	index	email_hash	orgyear	ctc	job_position	ctc_updated_year	company_hash_cleaned	Years_of_Experience	mean	median	max	
0	0	6de0a4417d18ab14334c3f43397fc13b30c35149d70c05...	2016	1100000	Other	2020	atrgxnnt xzaxv	8	1.100000e+06	1100000.0	1100000	1100000
1	1	b0aaf1ac138b53cb6e039ba2c3d6604a250d02d5145c10...	2018	449999	FullStack Engineer	2019	qtrxvzwt xzegwgbb rxbxnta	6	7.742856e+05	750000.0	1200000	449999
2	2	4860c670bcd48fb96c02a4b0ae3608ae6fdd98176112e9...	2015	2000000	Backend Engineer	2020	ojzwnvwnxw vx	9	2.000000e+06	2000000.0	2000000	2000000
3	3	effdede7a2e7c2af664c8a31d9346385016128d66bbc58...	2017	700000	Backend Engineer	2019	ngpgutaxv	7	1.436154e+06	1210000.0	3160000	700000
4	4	6ff54e709262f55cb999a1c1db8436cb2055d8f79ab520...	2017	1400000	FullStack Engineer	2019	qxen sqghu	7	1.400000e+06	1400000.0	1400000	1400000



Start coding or [generate](#) with AI.


```
# Standardize the numeric columns for clustering
scaler = StandardScaler()
df_scaled = scaler.fit_transform(df_cleaned[['ctc', 'Years_of_Experience']])
```

```
# Elbow method to find the optimal number of clusters
sse = []
for k in range(1, 11):
    kmeans = KMeans(n_clusters=k)
```




```
kmeans.fit(df_scaled)
sse.append(kmeans.inertia_)

# Plot the Elbow graph
plt.plot(range(1, 11), sse)
plt.title('Elbow Method')
plt.xlabel('Number of clusters')
plt.ylabel('SSE')
plt.show()
```




```
# Applying KMeans
kmeans = KMeans(n_clusters=3)
df_cleaned['Cluster'] = kmeans.fit_predict(df_scaled)

# Visualize the clusters
sns.scatterplot(x='Years_of_Experience', y='ctc', hue='Cluster', data=df_cleaned)
plt.show()
```




```
# Applying KMeans
kmeans = KMeans(n_clusters=2)
df_cleaned['Cluster'] = kmeans.fit_predict(df_scaled)

# Visualize the clusters
sns.scatterplot(x='Years_of_Experience', y='ctc', hue='Cluster', data=df_cleaned)
plt.show()
```




```
# Applying KMeans
kmeans = KMeans(n_clusters=4)
df_cleaned['Cluster'] = kmeans.fit_predict(df_scaled)

# Visualize the clusters
sns.scatterplot(x='Years_of_Experience', y='ctc', hue='Cluster', data=df_cleaned)
plt.show()
```




```
# Applying KMeans
kmeans = KMeans(n_clusters=5)
df_cleaned['Cluster'] = kmeans.fit_predict(df_scaled)

# Visualize the clusters
sns.scatterplot(x='Years_of_Experience', y='ctc', hue='Cluster', data=df_cleaned)
plt.show()
```



```
# Applying KMeans
kmeans = KMeans(n_clusters=6)
df_cleaned['Cluster'] = kmeans.fit_predict(df_scaled)

# Visualize the clusters
sns.scatterplot(x='Years_of_Experience', y='ctc', hue='Cluster', data=df_cleaned)
plt.show()
```



```
# Check the summary statistics of orgyear
print(df_cleaned['orgyear'].describe())
```



```
count    205665.000000
mean      2015.117584
std        4.228364
min       1976.000000
25%       2013.000000
50%       2016.000000
75%       2018.000000
max       2024.000000
Name: orgyear, dtype: float64
```

```
# Check for any unusual values (e.g., future years or very old years)
print(df_cleaned['orgyear'].value_counts().sort_index())
```



```
orgyear
1976      1
1977      1
1979      1
1981      1
1982      4
1984      3
1985      5
1986      8
1987      6
1988     10
1989     22
1990     38
1991     79
1992     47
```

```
1993      74
1994      65
1995      94
1996     134
1997     234
1998     279
1999     340
2000     495
2001     713
2002     685
2003    1018
2004    1455
2005    1873
2006    2075
2007    2257
2008    2728
2009    3777
2010    5751
2011    7970
2012   10493
2013   12351
2014   16696
2015   20610
2016   23043
2017   23239
2018   25256
2019   23427
2020   13431
2021    3670
2022     911
2023     252
2024      43
Name: count, dtype: int64
```

```
# Check the summary statistics of orgyear
print(df_cleaned['Years_of_Experience'].describe())
```

```
count    205665.000000
mean         8.882416
std         4.228364
min         0.000000
25%         6.000000
50%         8.000000
75%        11.000000
max        48.000000
Name: Years_of_Experience, dtype: float64
```

```
# Check for any unusual values (e.g., future years or very old years)
print(df_cleaned['Years_of_Experience'].value_counts().sort_index())
```

```
Years_of_Experience
0          43
1         252
2         911
3        3670
4       13431
5       23427
6       25256
```

```
7    23239
8    23043
9    20610
10   16696
11   12351
12   10493
13    7970
14    5751
15    3777
16    2728
17    2257
18    2075
19    1873
20    1455
21    1018
22     685
23     713
24     495
25     340
26     279
27     234
28     134
29      94
30      65
31      74
32      47
33      79
34      38
35      22
36      10
37       6
38       8
39       5
40       3
42       4
43       1
45       1
47       1
48       1
Name: count, dtype: int64
```

```
# # Perform hierarchical clustering
# Z = linkage(df_scaled, 'ward')
# dendrogram(Z)
# plt.show()
```

Start coding or [generate](#) with AI.

```
type(df1)
```